

Environmental News NLP Dataset Summarization

Project Report (Milestones 1, 2, and 3)

Noor Magdy (16005059)

Mohamed Wael (19011272)

Contents

1 Milestone 1: Data Understanding and Exploration	2
1.1 Introduction and Motivation	2
1.2 Literature Review	2
1.3 Data Analysis, Insights, and Limitations	3

1 Milestone 1: Data Understanding and Exploration

1.1 Introduction and Motivation

Natural Language Processing (NLP) tasks heavily depend on the quality and characteristics of the input data. Before applying any summarization techniques, it is essential to understand the structure, content, and limitations of the dataset being used. The goal of this project is to explore text summarization techniques on broadcast news data, evaluating both custom-built neural networks and state-of-the-art pre-trained transformer models.

The motivation behind this work stems from the increasing volume of televised and digital news. Providing accurate, concise summaries of environmental news allows policy-makers, researchers, and the general public to quickly grasp essential information without parsing through lengthy transcripts. Through basic data inspection, cleaning, exploratory data analysis, and subsequent modeling, this project establishes a solid foundation for comparing extractive and abstractive summarization methods.

1.2 Literature Review

Text summarization is a fundamental problem in Natural Language Processing, broadly categorized into two main approaches: extractive and abstractive summarization. Extractive summarization involves identifying and extracting the most salient sentences or phrases directly from the source text and concatenating them to form a summary. Traditional methods relied on statistical and graph-based algorithms, such as Term Frequency-Inverse Document Frequency (TF-IDF) and TextRank. These algorithms score sentences based on word frequencies or graph centrality. While extractive methods guarantee grammatical correctness (as they copy human-written text), they often suffer from poor coherence and an inability to condense information naturally.

The advent of Deep Learning introduced abstractive summarization, which aims to generate novel sentences that capture the core meaning of the source text, much like a human would. Early neural approaches relied on Sequence-to-Sequence (Seq2Seq) architectures powered by Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) networks. A standard Seq2Seq model utilizes an Encoder to compress the input text into a fixed-length context vector, and a Decoder to generate the summary token by token. However, fixed-length vectors struggle to retain information from long documents.

To resolve the bottleneck of fixed-length context vectors, the Attention Mechanism was introduced. Attention allows the decoder to dynamically "focus" on relevant parts of the input sequence during each step of the generation process. Additive attention and multiplicative attention significantly improved the fluency and accuracy of RNN-based abstractive summarizers.

Despite the success of RNNs with attention, they suffer from sequential processing constraints, making training slow and limiting their ability to model very long-range dependencies. This led to the paradigm shift brought by the Transformer architecture, which relies entirely on self-attention mechanisms and allows for massive parallelization.

Transformers paved the way for large pre-trained language models. For summarization, encoder-decoder transformers like BART (Bidirectional and Auto-Regressive Transformers) and T5 (Text-to-Text Transfer Transformer) represent the current state-of-the-art. BART is pre-trained by corrupting text with an arbitrary noising function and learning a model to reconstruct the original text. This denoising objective makes BART

exceptionally powerful for text generation and abstractive summarization tasks. In this project, we traverse this historical progression by first building an LSTM-based Seq2Seq model from scratch, followed by fine-tuning the pre-trained BART architecture.

1.3 Data Analysis, Insights, and Limitations

The dataset used in this project is the *Environmental News NLP Dataset*, obtained from Kaggle. It consists of short news snippets extracted from television news programs (CNN, MSNBC, BBC News, FOX News) related to environmental and public policy topics.

Data Overview & Insights: The dataset was downloaded into the Google Colab environment and combined from 418 separate CSV files, resulting in a single DataFrame of 94,858 records. Basic exploratory data analysis revealed the following:

- **Snippet Length Distribution:** The dataset contains snippets with an average length of approximately 41 words. The distribution is heavily concentrated between 30 and 50 words. The shortest snippet contains 2 words, while the longest reaches 194 words.
- **Program Consistency:** Boxplot analysis indicates that despite originating from different TV programs, the median snippet length remains consistent at around 40 words across all shows. Outliers are present in shows like BBC News and Hardball, indicating occasional extended commentary.
- **Temporal Stability:** A time-series analysis of average snippet lengths from 2009 to 2020 shows stable trends with only minor fluctuations, suggesting a standardized editorial style in broadcast news over the decade.

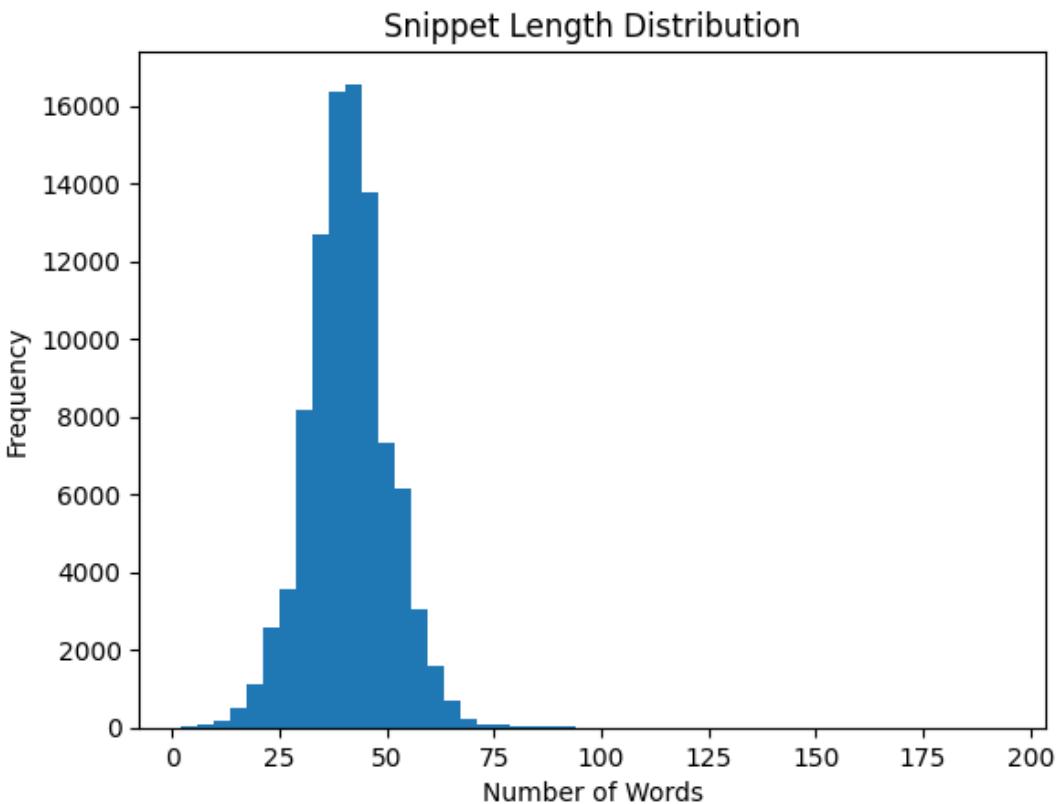


Figure 1: Exploratory Data Analysis showing snippet length distribution and program consistency.

Limitations: The primary limitation of this dataset is the absence of ground-truth reference summaries (labels). Because the texts are already relatively short (avg. 41 words), the potential for strong compression is limited. Furthermore, training supervised neural networks requires target labels. To overcome this in subsequent milestones, we utilize a frequency-based extractive heuristic to generate "silver standard" labels for training our abstractive models.